

Homework 2

Giulio Bagnoli

December 2019

Abstract: The paper presents an introduction to CNN and later a practical application of AlexNet and VGG11. The accuracy of these two networks calculated on the classification of the images contained in the Caltech-101 dataset will be analyzed. In particular, the work focuses on some techniques to obtain better and more robust networks such as:

- Transfer Learning.
- Data Augmentation.

1 Convolutional Neural Network

A Convolutional neural network (CNN) is a deep neural network that has one or more convolutional layers and it is used mainly for image processing, classification, segmentation and also for other auto correlated data.

CNNs have two components:

- **Feature extraction part:** in this part, the network will perform a series of convolutions and pooling operations during which the features are detected.
- **The Classification part:** here, the fully connected layers will serve as a classifier on top of these extracted features. They will assign a probability for the object on the image.

Convolutional networks were inspired by biological processes in that the connectivity pattern between neurons resembles the organization of the animal visual cortex. Individual neurons respond to stimuli only in a restricted region of the visual field known as the receptive field. The receptive fields of different neurons partially overlap such that they




Figure 1: AlexNet architecture scheme.

cover the entire visual field[1]. Therefore the network, rather than looking at an entire image at once to find certain features, divides the image in smaller portions from which it extracts the features. To do this, operations called convulsions are performed, a convolution is essentially sliding a filter over the input.

A CNN typically consists of 4 types of layers:

- Convolution Layer
- Pooling Layer
- Activation Function Layer
- Fully Connected Layer

1.1 Convolution Layer

The convolutional layer is the core building block of a CNN and it is always the first layer. The aim of this layer it is to extract the features from the input image that it is represented by *depth x height x width* array of pixel values where the dimension depending on the resolution and size of the image. Conventionally, the first ConvLayer is responsible for capturing the Low-Level features such as edges,

color, gradient orientation, etc. With added layers, the architecture adapts to the High-Level features as well.

In the context of a convolutional neural network, a convolution is a linear operation that involves the multiplication of a set of weights, called a filter or a kernel, with the input. The filter, that has a smaller height and width than the input data but has the same depth, is applied through calculation of the dot product between a filter-sized patch of the input and itself. The Bias term is added to the result of this product.

$$f(x, W) = W \cdot x + bias \quad (1)$$

Using a filter smaller than the input is intentional as it allows the same filter to be multiplied by the input array many times at different points on the input. Specifically, the filter is applied systematically to each overlapping part or filter-sized patch of the input data, left to right, top to bottom.




Figure 2: Movement of the Kernel.

This systematic application of the same filter across an image is a powerful idea. If the filter is designed to detect a specific type of feature in the input, then the application of that filter systematically across the entire input image allows the filter an opportunity to discover that feature anywhere in the image. This capability is commonly referred to as *translation invariance*.

As the filter is applied multiple times to the input array, the result is a two-dimensional array of output values that represent a filtering of the input. As such, the two-dimensional output array from this operation is called a *feature map*.

The dimensions of the feature map will be:

- $H_{out} = \frac{(H_{in} - H_K)}{S} + 1$




Figure 3: Filter application on the image.

- $W_{out} = \frac{(W_{in} - W_K)}{S} + 1$
- D_{out} = number of filters apply on the convolutional layer.

Where the subscript *in* represents the dimensions of the input image, the subscript *K* represents the dimensions of the filter and *S* represents the stride: the number of positions to which the filter is moved between one application and another.

Because the size of the feature map is always smaller than the input, something must be done to prevent our feature map from too shrinking. A layer of zero-value pixels is added to surround the input with zeros, so that our feature map will not shrink. In addition to keeping the spatial size constant after performing convolution, padding also improves performance and makes sure the kernel and stride size will fit in the input. If padding is applied the size of the padding used must be added two times to the input size in the output size calculation.

1.2 Pooling Layer

After a convolution layer, it is common to add a pooling layer between CNN layers. The pooling layer operates over each activation map independently and its function is to continuously reduce the dimensionality to reduce the number of parameters and computation in the network. This shortens the training time and controls overfitting. The pooling operation provides another form of translation invariance.

The most frequent type of pooling is max pooling, which partitions the input image into a set of rectangles and, for each sub-region, outputs the maximum. These sub-region sizes need to be specified beforehand.




Figure 4: Application of Maxpool.

This decreases the feature map size without introduce parameters while at the same time keeping the significant information.

Other pooling layer functions can be:

- Average pooling.
- L₂-norm pooling.
- Rol-pooling.

The dimension of the outputs can be calculated with the formulas of the ConvLayer except for the depth, which does not change.

1.3 Activation Function Layer

The activation function decides whether to activate the components of the output of the previous level. The purpose of this level is to introduce non-linearity in the network. An important feature of a activation function is that it should be differentiable. We need it to be this way so as to perform backpropagation optimization strategy while propogating backwards in the network to compute gradients of error(loss) with respect to Weights and then accordingly optimize weights using gradient descend or any other Optimization technique to reduce error.

The most used activation function is *ReLU*:

$$R(x) = \max(0, x) \quad (2)$$

One of the greatest advantage ReLU has over other activation functions is that it does not activate all neurons at the same time and it is less computationally expensive than *tanh* and *sigmoid* because it involves simpler mathematical operations. ReLU does not saturate at the positive region but it is saturated at the negative region, meaning that




Figure 5: ReLU.

the gradient at that region is zero. With the gradient equal to zero, during backpropagation all the weights will not be updated. To fix this, it can be used *Leaky ReLU*, which is similar to the ReLU with the only difference that in the negative region, instead of setting the value to zero (x), it sets it to $\alpha \cdot x$. Also, ReLU functions are not zero-centered. This means that for it to get to its optimal point, it will have to use a zig-zag path which may be longer.

1.4 Fully Connected Layer

Finally, after several convolutions and max pooling, the high-level reasoning in the neural network is done via fully connected (FC) layers. The objective of a FC layer is to take the results of the convolution/pooling process and use them to classify the image into a label, in a classification example.

Neurons in a FC layer have connections to all activations in the previous layer, as seen in regular (non-convolutional) artificial neural networks. Their activations can thus be computed with matrix multiplication followed by a bias offset. To do that the output of convolution/pooling is flattened into a single vector of values.

2 Training CNN

The overall training process of the Convolution Network may be summarized in the algorithm 2.

Where *step.size* indicates after how many epochs η is decreased by a γ factor, which has been set to 0.1 during the experiment.

Algorithm 1 Pseudocode of network training.

Data: *NUM_EPOCH*, *step_size*

Initialization of parameters with random values.

```

for epoch=1 To NUM_EPOCH do
  for batch in train dataloader do
    Performs forward pass.
    Calculates the loss.
    Perform backward pass.
  end
  if mod(epoch, step_size) == 0 then
    | Update  $\eta$ 
  end
end

```

Forward Pass: the image is passed over the network, the ConvLayer and Pooling layers act as feature extractors while FC layer acts as a classifier.

Calculates the Loss: once the network has determined the class of the input image the Loss is calculated. In the experiment carried out a particular Loss called *CrossEntropyLoss* is used which is expressed with:

$$\text{loss}(x, \text{class}) = \left(\log \frac{\exp(x[\text{class}])}{\sum_j \exp(x[j])} \right)^{(-1)} \quad (3)$$

Backward Pass: the backpropagation is applied, which computes the gradient of the loss function with respect to the weights of the network for a single input-output image by the chain rule. This makes it feasible to use gradient methods for training multi-layer networks and updating weights and biases to minimize loss. Commonly one uses gradient descent or variants such as stochastic gradient descent (SGD). A variant of the latter was used to conduct the experimentation.

2.1 Stochastic Gradient Descent + Momentum

Stochastic gradient descent is an optimization algorithm that estimates the error gradient for the current state of the model then updates the weights of the model using the back-propagation of errors algorithm. The amount that the weights are updated during training is referred to as the learning rate (η).

$$w_{t+1} = w_t - \eta \nabla L(w_t) \quad (4)$$

Specifically, η is a configurable hyperparameter used in the training of neural networks that has a small positive value, often in the range between 0.0 and 1.0.

The learning rate controls how quickly the model is adapted to the problem. Smaller learning rates require more training epochs given the smaller changes made to the weights each update, whereas larger learning rates result in rapid changes and require fewer training epochs. A learning rate that is too large can cause the model to converge too quickly to a suboptimal solution, whereas a learning rate that is too small can cause the process to get stuck. η may be the most important hyperparameter for the model.

SGD presents some problems if there are local minimums or saddles. For this reason the version of SGD + Momentum has been used. It allows us to better escape from these points, converging more quickly to a global solution. To do this the weights are updated differently:

$$v_{t+1} = \rho v_t + \nabla L(w_t) \quad (5)$$

$$w_{t+1} = w_t - \eta v_{t+1} \quad (6)$$

In the experiment rho value was set to 0.9.

3 Transfer Learning (TL)

Transfer learning is a machine learning method where a model developed for a task is reused as the starting point for a model on a second task. So TL is a research problem that focuses on storing knowledge gained while solving one problem and applying it to a different but related one. Transfer learning is an optimization that allows rapid progress or improved performance when modeling the second task. Another motivation to use TL is that, especially considering the context of deep learning, most models which solve complex problems need a whole lot of data, and getting vast amounts of labeled data for supervised models can be really difficult. Through the use of TL it is possible to recover the knowledge from data that do not belong at the task that it has to be performed and obtain good result also begin from a small amount of data.




Figure 6: AlexNet architecture scheme.

The two most popular strategies for deep transfer learning are:

- **Pre-trained Models as Feature Extractors:** Deep learning models are layered architectures that learn different features at different layers (ConvLayer). These layers are then finally connected to a last layer (usually a fully connected layer, in the case of supervised learning) to get the final output. This layered architecture allows to utilize a pre-trained network without its final layer as a fixed feature extractor for other tasks.

For instance, if AlexNet is used without its final classification layer, it will help transform images from a new domain task into a 4096-dimensional vector based on its hidden states, thus enabling to extract features from a new domain task, utilizing the knowledge from a source-domain task.

- **Fine Tuning Pre-trained Models:** This is a more involved technique, where it is done not just replace the final layer (for classification/regression), but also it is retrained some of the previous layers. Deep neural networks are highly configurable architectures with various hyperparameters. As discussed earlier, the initial layers have been seen to capture generic features, while the later ones focus more on the specific task at hand

Using this insight, certain layers may be frozen while retraining, or fine-tune the rest of them to suit needs. In this case, knowledge is utilized in terms of the overall architecture of the network and use its states as the starting point for retraining step. This helps achieve better performance with less training time.

3.1 Freezed Layers

Freezing a layer prevents its weights from being modified. This technique is often used in transfer learning, where the base model is frozen. If some layer are frozen, the backward pass to that layer can be completely avoided, resulting in a significant speed boost. On the other hand, you still need to train the model, so if you freeze it too early, it will give inaccurate predictions. This technique is to cut down on the computational time for training while losing not much on the accuracy side.

4 Data Augmentation

Data Augmentation is a technique that allows to solve some problems related to the lack of a sufficient amount of data. In particular Data Augmentation is based on the modification of the input images of the model, this is possible since a CNN is able to classify an object even if it is placed in different orientations. This property is called invariance. More specifically, a CNN can be invariant to translation, viewpoint, size or illumination, or a combination of the these.

Data Augmentation is done before the images are given to the CNN, this it can be done in two different ways: one option is to perform all the necessary transformations beforehand, essentially increasing the size of your dataset. The other option is to perform these transformations on a mini-batch, just before feeding it to your machine learning model.

The first option is known as **offline augmentation**. This method is preferred for relatively smaller datasets, since the size of the dataset will increase by a factor equal to the number of transformations performed.

The second option is known as **online augmentation**, or augmentation on the fly. This method is preferred for larger data sets, since the explosive size increase cannot be allowed. This procedure was used in the experiment.

Data Augmentation cannot overcome all biases present in a small dataset. For example, in a traffic signs classification task, if there are only signs of obligation and no case of stop signals, no augmentation method will create a stop signals. However, several forms of biases such as lighting, occlusion, scale, background, and many more are preventable

or at least dramatically lessened with Data Augmentation. Overfitting is generally not as much of an issue with access to big data. Data Augmentation prevents overfitting by modifying limited datasets to possess the characteristics of big data.

The most common transformations are:

- *Flip*: it is possible flip images horizontally or vertically.
- *Rotation*: one key thing to note about this operation is that image dimensions may not be preserved after rotation
- *Scale*: the image can be scaled outward or inward.
- *Random Cropping*: it just randomly sample a section from the original image and then resize this section to the original size.

5 Early Stopping

In machine learning, early stopping is a form of regularization used to avoid overfitting when training a model with an iterative method, such as gradient descent. Such methods update the model so as to make it better fit the training data with each iteration. Therefore if the model is trained too little it will commit underfitting on the train and the test sets. Whereas if it trained too much it will commit overfitting on the training set and have poor performance on the test set.

The challenge is to train the network long enough that it is capable of learning the mapping from inputs to outputs, but not training the model so long that it overfits the training data. A compromise is to train on the training set but to stop training at the point when performance on a validation set starts to degrade.

To understand when the performance of the model on the validation set is decreasing due to overfitting, various techniques have been studied, explained in [2]. In particular, in the experiment, referring to the terminology of the article mentioned above, an early stop criterion of the UP type with parameters ($K = 5$, $S = 5$) was developed.

6 AlexNet

AlexNet, which is the network used for experiments, competed in the ImageNet Large Scale Visual Recognition Challenge on September 30, 2012. The network achieved a top-5 error of 15.3%, more than 10.8 percentage points lower than that of the runner up.

The input to AlexNet is an RGB image of size 224x224. This means all images used need to be of size 224x224. If the input image does not have these dimensions, it needs to be converted to 224x224 before using it for training the network.




Figure 7: AlexNet architecture scheme.

The figure 8 represents the architecture of the network, the list of parameters of the various levels is in this GitHub repository.

7 Description of the project

In the project AlexNet was trained on the Caltech-101 dataset, present in this GitHub repository. The following aspects were analyzed:

- Variations of the accuracy of the network with η and *step_size* changes.
- Variations of the accuracy of the network with η and *step_size* changes with the pre-trained network on the dataset ImageNet.
- Study on the effects of the application of Data Augmentation.

7.1 Code Overview

Caltech-101 is a dataset containing 102 classes of different images, 101 of these represent real objects while the last represents only backgrounds. For this

reason it was necessary to create a class, which inherits from *VisionDataset*, which during the loading of the dataset excluded the background category. This class, called Caltech, is present in the *Utility* cell. The following methods are also present in this cell:

- *prepareDataloader*: it is used to create the dataloaders to scroll through the various sets (train, validation, test and train + validation) during the training and calculation phases of accuracy. The first point of the homework (*Data Preparation*) is entirely developed by this method.

The method has as input an object of the *transform* type that modifies the images normalizing and resizing them in the first phases, while in the case of Data Augmentation it allows to apply the various transformations.

- *stoppingCriteria*: it defines a criterion to stop the training of the network before the maximum number of epochs is reached. It allows therefore to have faster trainings in the cases in which the accuracy of the network is constant or is decreasing.
- *partialTrain*: it trains the net on the union of train and validation by freezing the convolutional or the FC layers and using as hyperparameters those passed in input. The choice of which part of the network to freeze also depends on the input. Once the network has been trained, the accuracy of the test set is calculated.
- *tuneHyperparameters*: it is the method in which various hyperparameters are tested. Inside it contains a cycle in which the network is trained with the various combinations of $(\eta, step_size)$. For each epoch the accuracy on the validation is calculated and, once the major one is determined, the network is retrained on the union of the train and the validation and its accuracy on the test set is calculated. This last training phase is carried out using the set of hyperparameters $(\eta, step_size, num_epoch)$ which gave the best accuracy in the previous phase. It is noteworthy that the number of epochs with which the

network is trained the second time is not necessarily the maximum number of epochs defined but it is the one for which the best accuracy was obtained. In the case in which the loss value, calculated in the training phase, is *NaN* this phase is interrupted and the pair of hyperparameters is discarded.

The method has a network as input and the necessary dataloaders to scroll through the various sets, in this way it can be used both in the initial phase, where the network is trained from scratch, and in the subsequent phases, in which Learning Transfer is carried out.

Algorithm 2 Pseudocode of *tuneHyperparameters*.

Data: *NUM_EPOCH*, *step_size*, η
accuracy_max = 0
best_hyperparameters = [0,0]
for each couple of hyperparameters **do**
 for epoch=1 To *NUM_EPOCH* **do**
 Train the network.
 Calculation of accuracy on the validation set.
 if accuracy > accuracy_max **then**
 update best_hyperparameters
 update accuracy_max
 end
 end
end

end

Train the network with best_hyperparameters.

Calculation of accuracy on the test set.

Since AlexNet was designed to classify 1000 images while the dataset used has only 101 all the methods in which the network is trained modify the number of outputs of the last FC level by setting it to 101.

The couple of hyperparameters $(\eta, step_size)$ tested are showed in table 1 and they are present in the code inside the *Set Argument* cell. The values for these hyperparameters have been chosen so that the high η fall rapidly, while the low ones vary very slowly.

η	<i>step_size</i>
0.1	20
0.01	30
0.01	45
0.001	60

Table 1: Values of η , *step_size*.

Set Argument cell also contains the remaining parameters to train the network as:

- *NUM_EPOCH*: is the maximum number of epochs for which the network is trained, it is set to 200.
- *BATCH_SIZE*: is the dimension of each batch, it is set to 256 for AlexNet while for VGG11 it is set to 32.
- *MOMENTUM*, *WEIGHT_DECAY*, *GAMMA*: are parameters necessary for the calculation of the SGD algorithm, their values are (0.9, $5 * 10^{-5}$, 0.1), as mentioned above.
- *HP*, *strip_len*: are parameters necessary for the calculation of the early stop algorithm, they represent the maximum number of consecutively controlled strips and the length of each strip. Both are set to 5.
- *train*, *TL*, *data_augmentation*(*_1/_2/_3*) *transform*: are the objects that will be passed to the *prepareDataloader* method. All carry out the transformations necessary to bring the dimensions of the images from those original to those permitted by the network. Since *TL_transform* and *data_augmentation_transforms* are used to generate dataloaders that will scroll through images for pre-trained networks, these transforms normalize images with different values than *train_transform*. Furthermore *data_augmentation_transform* is used for real-time data augmentation so it also implements a set of transformations.

7.2 Training from scratch

After creating both the dataloaders, using the *PrepDataloader* method to which *train_transform* has been passed, and the untrained network, which is provided by Pytorch in both version trained and

untrained, the *tuneHyperparameters* method is executed and its whose results are shown below.

For the couple of hyperparameters ($\eta = 0.1$, *step_size* = 20) the best accuracy (16.25) is reached at the epoch 6. The train is stopped to the epoch 8 because the value of the loss is *NaN*, this happen because the high value of η make diverge the loss.

For the second couple of hyperparameters ($\eta = 0.01$, *step_size* = 30) the best accuracy (47.61) is reached at the epoch 92. The lowest η does not lead to the divergence of the loss allowing the network to train for longer. The training phase ends at the iteration 185 due to the stop criterion. In the last 100 epochs it is not possible to notice big changes in the accuracy calculated on the validation because, due to the very low *stepsize*, the weights of the net are updated of very small quantities. The early stop criterion does not stop the training phase beforehand as the accuracy continues to improve slightly, until the last 25 epochs.

For the third pair of hyperparameters ($\eta = 0.01$, *step_size* = 45) the best accuracy is (53.63) is reached at the epoch 90. Having a higher step size η increases more slowly allowing more significant weight changes. As in the previous case it is possible to find how the accuracy remains more or less unchanged once η assumes a value of 10^{-4} .

For the fourth pair of hyperparameters ($\eta = 0.001$, *step_size* = 60) the best accuracy is (9.19) is reached at the epoch 3. After which the accuracy remains constant for the next 25 epochs and the early stop criterion stops the training.

The best value obtained for accuracy was that of the third pair of hyperparameters, therefore the network was retrained with these values for 90 epochs on the union of the train and the validation. Later, the accuracy of the 50.92 test set was calculated.

7.3 Transfer Learning

Since Caltech-101 is a relatively small dataset on which to train a CNN, it was decided to implement transfer learning as an attempt to improve the accuracy of the network. The TL mode used was *Fine Tuning Pre-trained Models* as in the entire pre-trained AlexNet it was trained on Caltech-101. To do this, the data loaders were created using the *TL_transform* transform object.

The results, obtained by executing *tuneHyperparameters* are showed below. As for the previous point, the pair of hyperparameters ($\eta = 0.1$, *step_size* = 20) caused the loss on the train to diverge within the first epochs, thus interrupting the training of the network. On the other hand the other hyperparameter pairs have obtained significantly higher accuracy than the previous case. The results in detail are shown in tables 2

η , <i>step_size</i>	Max Acc and Epoch	Epoch of exit
(0.01, 30)	(86.30%, 52)	120
(0.01, 45)	(87.21%, 90)	115
(0.001, 60)	(85.13%, 113)	180

Table 2: TL results.

As expected, the pre-trained model is able to recognize the general features of the image, obtaining from the first iterations a highest accuracy with respect to all epochs of previous point. Accuracy continues to improve in later epochs when the model learns the high-level features of the images contained in Caltech-101.

The result of the trained network, both on the train and on the validation, with the hyperparameters (0.01, 45) is 85.00%.

Subsequently the network was retrained two more times, using the best pair of hyper-parameters obtained in the previous step, freezing in the first training the convolutional layers and in the second the fully connected layers. The accuracy calculated on the test in the two cases is very different, in the first case in fact the network reaches a value of 88.52%, while in the second case only 72.42%.

These phase of training took much less time than the previous ones, this is due to the fact that, freezing half of the levels, the update phase concerns fewer parameters, and is therefore faster.

7.4 Data Augmentation

To further increase the volume of the datasets some Data Augmentation techniques have been applied.

The objects *data_augmentation_transform* include within themselves the following transformations:

- *RandomCrop*: Crop the given PIL Image at a random location.
- *RandomHorizontalFlip*: Horizontally flip the given PIL Image randomly with 50% probability.
- *Grayscale*: Convert image to grayscale. The number of the channels of the output image is set to 3.
- *RandomVerticalFlip*: Vertically flip the given PIL Image randomly with 50% probability.
- *RandomRotation*: Rotate the image by angle between -90 and 90 degrees.
- *ColorJitter*: Randomly change the brightness, contrast and saturation of an image.

The following combinations of transformations have been tested:

- RandomCrop, Grayscale, RandomHorizontalFlip.
- Grayscale, RandomCrop, RandomVerticalFlip.
- RandomCrop, RandomHorizontalFlip, ColorJitter.

For each combination of transformations the *tuneHyperparameters* method was performed. The best triad of transformations is: *RandomCrop*, *RandomHorizontalFlip*, *ColorJitter* with. This set of transformations allows the network to reach an maximum accuracy of 74.59% on the test set. The other two sets of transformations only reach the 72.21% and the 68.89% of accuracy.

While before there was a strong similarity between the images of the train and the test set now there is no more, this leads to a general decrease in network performance compared to when only transfer learning was implemented. However, techniques such as Data Augmentation allow the network to commit overfitting in more advanced stages of training, making the network able to better classify data that is different from the train set.

7.5 Beyond AlexNet

To further study the behavior of CNN some of the previous experiments have been conducted on VGG11. VGG is a CNN architecture trained by Oxford's Visual Geometry Group. VGG11 is the smallest model of the VGG architecture implemented by Pytorch.

VGG's architecture instead of using large filter sizes with large strides, such as in AlexNet, uses small (3x3) filter sizes with stride 1 throughout the whole net. Moreover all the inputs to ConvLayers are appropriately padded such that after the convolution step the spatial dimension of the output is same as the input.

Only the pooling layers are responsible for changing the spatial dimensions. All max-pooling layers use a (2x2) filter with stride 2. This results in the reduction of spatial dimensions by a factor of 2. A ReLU layer is applied to every output of the ConvLayers and the FC layers.




Figure 8: Comparison between AlexNet and VGG architecture schemes.

A stack of two (3x3) ConvLayers has an effective receptive field of one (5x5) ConvLayer; and three such ConvLayers have an effective receptive field of one (7x7) ConvLayer. As a result, by using multiple small filter we are not losing any information, in terms of receptive field area.

Comparison between three (3x3) Conv Layers and a (7x7) Conv Layer.

- Both have the same receptive field.
- As each ConvLayer is associated with a ReLU layer if more than one small filter is used, a greater non-linearity is introduced in the network. This makes the decision function more discriminative and hence learn more complex

features. This also allows the model to create a better mapping from the images to the labels.

- The number of parameters decrease using the stack of (3x3) ConvLayers. Assuming both the input and output have C channels, the stack of (3x3) is parametrised by $3(3C) = 27C$ weights whereas the (7x7) will require $1(7C) = 49C$ weights. Using (3x3) ConvLayers decreases the size of the model on memory and also acts as a sort of regularisation, making the network less prone to overfitting.

So using multiple layers with small filters is more advantageous than using fewer ones with larger filters if a layer with an activation function is added to each convolutional layer.

As VGG11 is deeper than AlexNet, the Batch size has been reduced from 256 to 64. This leads to more weight updates by age and more training time.

Generally VGG11 performs better than AlexNet in all the tests involved, except those in which the combination of hyperparameters ($\eta = 0.1, step_size = 20$) has been used. This pair of hyperparameters leads even VGG11 to have a loss equal to NaN. The results of the various experiments are shown in Table 3. Noteworthy is the accuracy obtained by the pretrained network with the pair of hyperparameters ($\eta = 0.001, step_size = 60$), 90.7% in fact it is the best result obtained on the test in all the experiments performed.

Test Performed	Hyperameters ($\eta, step_size, epoch$)	Accuracy on test set
Trained from scratch	(0.01, 30, 55)	52.3%
TL	(0.001, 60, 147)	90.7%
Data Augmentation	(0.001, 60, 101)	86.76%

Table 3: Result obtained by VGG11.

To do Data Augmentation all three combinations of transformations are used. The higher accuracy

obtained from Data Augmentation is present in Table 3 and it is calculated with the last combination of transformations.

References

- [1] *Convolutional neural network*. URL: https://en.wikipedia.org/wiki/Convolutional_neural_network#cite_note-fukuneoscholar-5.
- [2] Lutz Prechelt. “Early Stopping — but when?” In: (). URL: http://page.mi.fu-berlin.de/prechelt/Biblio/stop_tricks1997.pdf.